


POSIX Threads in a nutshell

Paolo Burgio
paolo.burgio@unimore.it





What will we see

- › A mix of theory...
- › ..and practice / exercise
 - Don't miss it
- › Please, interrupt me



The POSIX IEEE standard

eng.wikipedia.org

POSIX Threads, usually referred to as Pthreads, is an execution model that exists independently from a language, as well as a parallel execution model. It allows a program to control multiple different flows of work that overlap in time.

- › Threading API
- › Single process
- › Shared memory space





The POSIX IEEE standard

- › Specifies an **operating system interface similar to most UNIX systems**
 - It extends the C language with primitives that allows the specification of the concurrency
- › POSIX distinguishes between the terms process and thread
 - "A **process** is an address space with one or more threads executing"
 - "A **thread** is a single flow of control within a process (a unit of execution)"
- › Every process has at least one thread
 - the `main()` (aka "**master**") thread; its termination ends the process
 - All the threads **share** the same address space, and have a **private** stack



Thread body

- › A (P)thread is identified by a C function, called body:

```
void *my_pthread_fn(void *arg)
{
    // Thread body
}
```

- › A thread starts with the first instruction of its body
- › The threads ends when the body function ends
 - It's not the only way a thread can die



Thread creation

- › Thread can be created using the primitive

pthread.h

```
typedef unsigned int pthread_t;  
  
int pthread_create ( pthread_t *ID,  
                    pthread_attr_t *attr,  
                    void *(*body)(void *),  
                    void * arg  
                    );
```

- › pthread_t is the type that contains the thread ID
- › pthread_attr_t is the type that contains the parameters of the thread
- › arg is the argument passed to the thread body when it starts



Thread attributes

- › Thread attributes specifies the characteristics of a thread
 - We won't see this; leave empty
- › Attributes must be initialized and destroyed - **always**

pthread.h

```
int pthread_attr_init(pthread_attr_t *attr);  
  
int pthread_attr_destroy(pthread_attr_t *attr);
```



Thread termination

- › A thread can terminate itself calling

pthread.h

```
void pthread_exit(void *retval);
```

- › When the thread body ends after the last “}”,
pthread_exit() is called implicitly
- › Exception: when main() terminates, exit() is called implicitly



Thread IDs

- › Each thread has a unique ID

pthread.h

```
pthread_t pthread_self(void);
```

- › The thread ID of the current thread can be obtained using

pthread.h

```
int pthread_equal( pthread_t thread1,  
                  pthread_t thread2 );
```

- › Two thread IDs can be compared using



Joining a thread

- › A thread can wait the termination of another thread using

pthread.h

```
int pthread_join ( pthread_t th,  
                  void **thread_return);
```

- › It gets the return value of the thread or `PTHREAD_CANCELED` if the thread has been killed
- › By default, every thread **must** be joined
 - The join frees all the internal resources
 - Stack, registers, and so on



Example

Let's
code!

- › Filename: `hello_pthreads_world.c`
- › The demo explains how to create a thread
 - the `main()` thread creates another thread (called `body()`)
 - the `body()` thread checks the thread ids using `pthread_equal()` and then ends
 - the `main()` thread joins the `body()` thread
- › When compiling under gcc & GNU/Linux, remember
 - the `-lpthread` option!
 - to add `#include "pthread.h"`

› Credits to PJ





Detached threads

- › A thread which does not need to be joined has to be declared as **detached**
- › 2 ways to have it:
 - While creating (in father thread) using `pthread_attr_setdetachstate()`
 - The thread itself can become detached calling in its body `pthread_detach()`
- › Joining a detached thread returns an **error**



Killing a thread

- › A thread can be killed calling

```
int pthread_cancel(pthread_t thread);
```

- › When a thread dies its data structures will be released
 - By the join primitive if the thread is joinable
 - Immediately if the thread is *detached*
 - Why?



PThread cancellation

- › Specifies how to react to a **kill** request
- › There are two different behaviors:
 - **deferred** cancellation
when a kill request arrives to a thread, the thread does not die. The thread will die only when it will execute a primitive that is a **cancellation point**. This is the default behavior of a thread.
 - **asynchronous** cancellation
when a kill request arrives to a thread, the thread dies. The programmer must ensure that all the application data structures are coherent.



Cancellation states and cleanups

- › The user can set the cancellation state of a thread using:

```
int pthread_setcancelstate(int state, int *oldstate);  
int pthread_setcanceltype(int type, int *oldtype);
```

- › The user can protect some regions providing destructors to be executed in case of cancellation

```
int pthread_cleanup_push(void (*routine)(void *),  
                        void *arg);  
int pthread_cleanup_pop(int execute);
```



Cancellation points

- › The **cancellation points** are primitives that can potentially **block** a thread
- › When called, if there is a kill request pending the thread will die
 - `void pthread_testcancel(void);`
 - `sem_wait`, `pthread_cond_wait`, `printf` and all the I/O primitives are **cancellation points**
 - `pthread_mutex_lock`, is **NOT** a cancellation point
 - › Why?
- › A complete list can be found into the POSIX Standard



Cleanup handlers

- › The **user** must **guarantee** that when a thread is killed, the application data remains coherent.
- › The user can **protect** the application code using **cleanup handlers**
 - A cleanup handler is a user function that *cleans up* the application data
 - They are called when the thread **ends** and when it is **killed**



Cleanup handlers (2)

```
void pthread_cleanup_push (void (*routine) (void *),  
                           void *arg);  
  
void pthread_cleanup_pop (int execute);
```

- They are pushed and popped as in a stack
- If `execute != 0` the cleanup handler is called when popped
- The cleanup handlers are called in LIFO order



How to run the examples

Let's
code!

› Download the Code/ folder from the course website

› Compile

```
$ gcc code.c -o code -lpthread
```

› Run (Unix/Linux)

```
$ ./code
```

› Run (Win/Cygwin)

```
$ ./code.exe
```

Useful links



› Course webpage

- https://hipert.unimore.it/people/paolob/pub/Calcolo_Parallelo/

› Course GitHub

- <https://github.com/HiPeRT/cp19/>



› My contacts

- paolo.burgio@unimore.it
- <http://hipert.mat.unimore.it/people/paolob/>

› PThreads

- <https://computing.llnl.gov/tutorials/pthreads/>
- <http://man7.org/linux/man-pages/man7/pthreads.7.html>

› A "small blog"

- <http://www.google.com>