

# Critical sections in OpenMP

---

Paolo Burgio  
paolo.burgio@unimore.it



# Outline

---

- › Expressing parallelism
  - Understanding parallel threads
- › ~~Memory~~ Data management
  - Data clauses
- › Synchronization
  - Barriers, locks, critical sections
- › Work partitioning
  - Loops, sections, single work, tasks...
- › Execution devices
  - Target



# OpenMP synchronization

---

- › OpenMP provides the following synchronization constructs:
  - barrier
  - flush
  - master
  - critical
  - atomic
  - taskwait
  - taskgroup
  - ordered
  - ...and OpenMP locks



# Exercise

---

Let's  
code!

- › Spawn a team of (many) parallel Threads
  - Each incrementing a shared variable
  - What do you see?



# Coherency problem

---

- › When multiple workers concurrently access the same shared data
  - Data races
  - Read stale data
  - ...
  
- › (Luckily, does not happen with read-only data)



# Data races, aka: the a++ problem

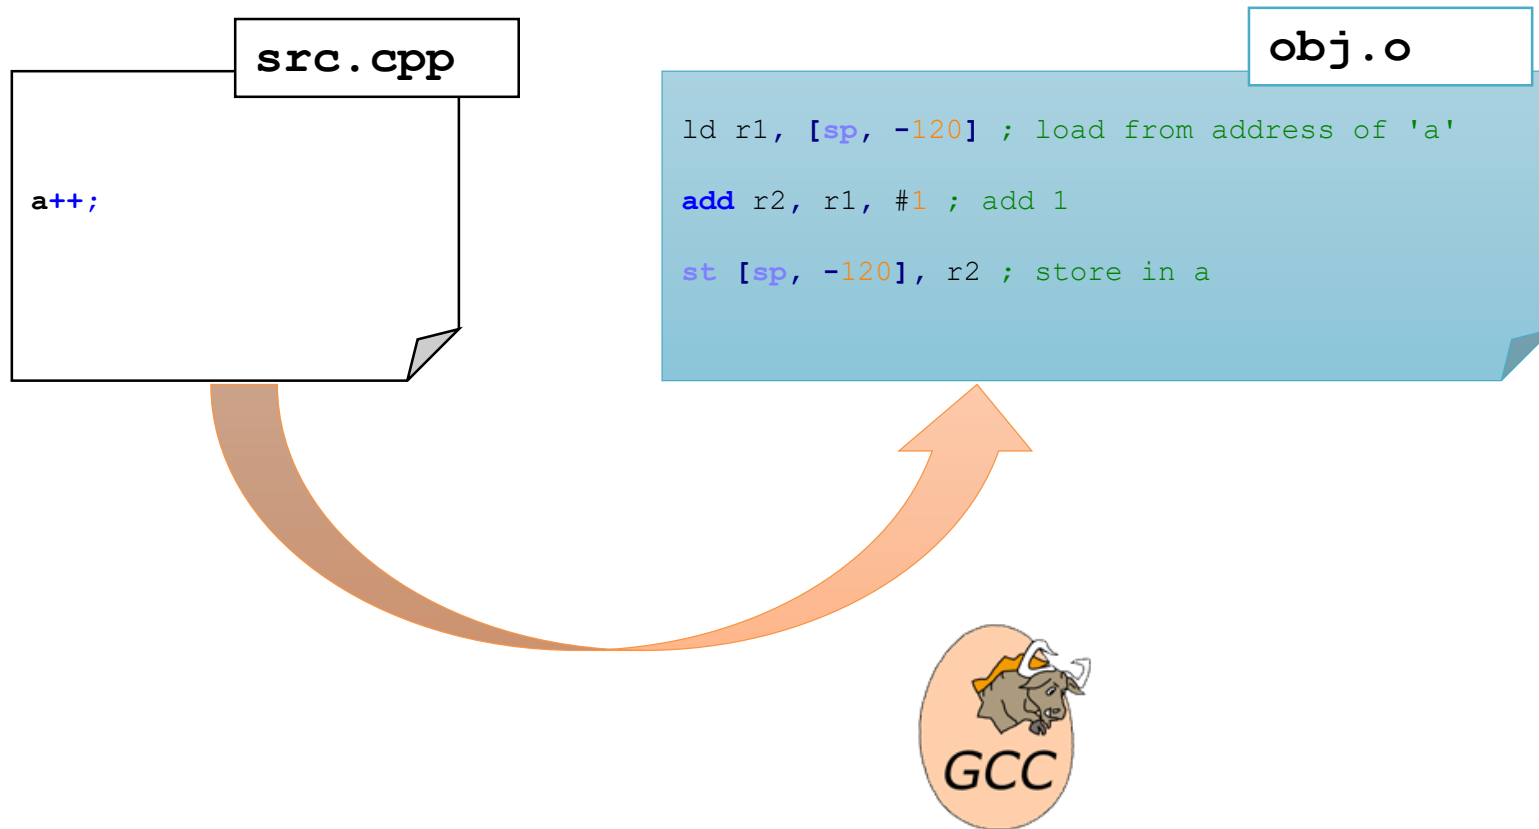
---

- › Intrinsic to the way computers are build
  - Leave out *Processor-in-memory*, ok?
- › A shared variable `a`
- › Incremented by two parallel threads
- › Interleaved execution
- › Eisen-bug vs Bohr-bug



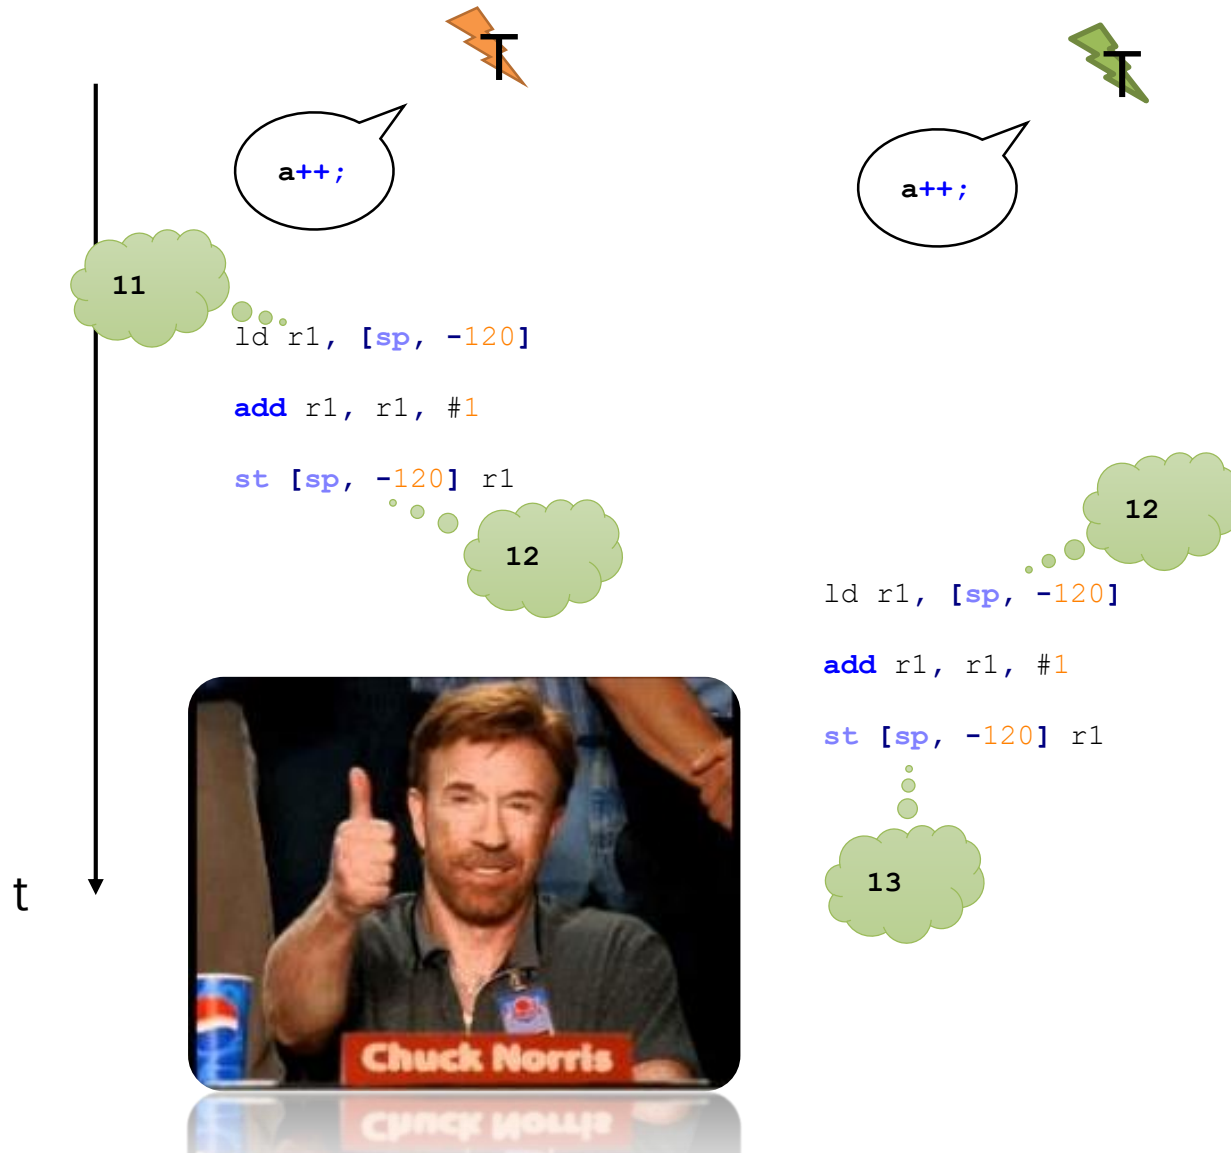
# The big problem

- › a++ is not an unique instruction






# When things go well









# OpenMP locks

---

- › Defined at the OpenMP runtime level
  - Symbols available in code including `omp.h` header
  
- › General-purpose locks
  1. Must be initialized
  2. Can be set
  3. Can be unset
  
- › Each lock can be in one of the following states
  1. Uninitialized
  2. Unlocked
  3. Locked





# Locking primitives

omp.h

```
/* Initialize an OpenMP lock */  
void omp_init_lock(omp_lock_t *lock);  
  
/* Ensure that an OpenMP lock is uninitialized */  
void omp_destroy_lock(omp_lock_t *lock);  
  
/* Set an OpenMP lock. The calling thread behaves  
   as if it was suspended until the lock can be set */  
void omp_set_lock(omp_lock_t *lock);  
  
/* Unset the OpenMP lock */  
void omp_unset_lock(omp_lock_t *lock);
```

- › The `omp_set_lock` has blocking semantic
- › `omp_set/unset_lock` are thread safe
- › Opaque data type `omp_lock_t`



# OMP locks: example

## › Locks **must** be

- Initialized
- Destroyed

## › Locks can be

- set
- unset
- tested

## › Very simple example

```
/** Do this only once!! */
/* Declare lock var */
omp_lock_t lock;
/* Init the lock */
omp_init_lock(&lock);

/* If another thread set the lock,
   I will wait */
omp_set_lock(&lock);

/* I can do my work being sure that no-
   one else is here */

/* unset the lock so that other threads
   can go */
omp_unset_lock(&lock);

/** Do this only once!! */
/* Destroy lock */
omp_destroy_lock(&lock);
```



# Exercise

---

Let's  
code!

- › Spawn a team of (many) parallel Threads
  - Each incrementing a shared variable
  - What do you see?
  
- › Protect the variable using OpenMP locks
  - What do you see?
  
- › Now, comment the call to `omp_unset_lock`
  - What do you see?



# Non-blocking lock set

omp.h

```
/* Set an OpenMP lock but do not suspend the execution of the thread.  
Returns TRUE if the lock was set */
```

```
int omp_test_lock(omp_lock_t *lock);
```

- › Extremely useful in some cases. Instead of blocking
  - we can do useful work
  - we can increment a counter (to profile lock usage)
  
- › Reproduce blocking set semantic using a loop
  - `while (!omp_test_lock(lock)) /* ... */;`



# The `omp_lock_t` type

`omp.h`

```
/* (1) Our implementation @UniBo (few years ago) */
typedef unsigned long omp_lock_t;

/* (2) ROSE compiler */
typedef void * omp_lock_t;


/* (3) GCC-OpenMP (aka Libgomp) */
typedef struct {
    unsigned char _x[@OMP_LOCK_SIZE@]
        __attribute__((__aligned__(@OMP_LOCK_ALIGN@)));
} omp_lock_t;
```

- › Opaque data type
- › Implementation-defined, it represents a lock type
  - Different implementations, different optimizations
- › C routines for OMP lock are thread-safe, and accept a pointer to an `omp_lock_t` type
  - (at least)

# Exercise

Let's  
code!

- › Modify the "PI Montecarlo" exercise
  - Replace the variable in the `reduction` clause with a `shared` variable
  - Protect it using an OpenMP lock







# Let's do more

---

- › Locks are extremely powerful
  - And low-level
- › We can use them to build complex semantics
  - Mutexes
  - Semaphores..
- › But they are a bit "cumbersome" to use
  - Need to initialize before, and release after
  - We can definitely do more!

pragma-level synchronization constructs



# The critical construct

```
#pragma omp critical [(name) [hint(hint-expression)] ] new-line  
    structured-block
```

Where *hint-expression* is an integer constant expression that evaluates to a valid lock hint

- › "Restricts the execution of the associated structured block to a single thread at a time"
  - The so-called Critical Section
- › Binding set: all threads everywhere (also in other teams/parregs)
- › Can associate it with a "hint"
  - `omp_lock_hint_t`
  - Also locks can
  - We won't see this



# The critical section

› From this...

```
/* Declare lock var */
omp_lock_t lock;
/* Init the lock */
omp_init_lock(&lock);

/* If another thread set the lock,
   I will wait */
omp_set_lock(&lock);

/* I can do my work being sure that no-
   one else is here */

/* unset the lock so that other threads
   can go */
omp_unset_lock(&lock);

/* Destroy lock */
omp_destroy_lock(&lock);
```

› ...to this

```
/* If another thread is in, I must wait */

#pragma omp critical
{
    /* _Critical Section_
       I can do my work being sure
       that no- one else is here */

}

/* Now, other threads can go */
```




# Exercise

---

Let's  
code!

- › Modify the "PI Montecarlo" exercise
  - Using critical section instead of locks




# The risk of sequentialization



## › Critical sections should be kept small as possible

- They force code portions sequentialization
- Harness performance





# Even more flexible

---



- › (Good) parallel programmers manage to keep critical sections small
  - Possibly, away from their code!
- › Most of the operations in a critical section are always the same!
  - "Are you really sure you can't do this using `reduction` semantics?"
  - Modify a shared variable
  - Enqueue/dequeue in a list, stack..
- › For single (C/C++) instruction we can definitely do better



# The atomic construct

```
#pragma omp atomic [seq_cst] new-line  
expression-stmt
```


- › The atomic construct ensures that a specific storage location is accessed atomically
  - We will see only its simplest form
  - Applies to a single instruction, not to a structured block..
  
- › Binding set: all threads everywhere (also in other teams/parregs)
  
- › The `seq_cst` clause forces the atomically performed operation to include an implicit `flush` operation without a list
  - Enforces memory consistency
  - Does not avoid data races!!



# Exercise

Let's  
code!

- › Modify the "PI Montecarlo" exercise
  - Implementing the critical section with the `atomic` construct
  - (If possible)







# How to run the examples

---

Let's  
code!

› Download the Code/ folder from the course website

› Compile

› `$ gcc -fopenmp code.c -o code`

› Run (Unix/Linux)

`$ ./code`

› Run (Win/Cygwin)

`$ ./code.exe`

# References

---



- › "Calcolo parallelo" website
  - [http://hipert.unimore.it/people/paolob/pub/Calcolo\\_Parallelo/](http://hipert.unimore.it/people/paolob/pub/Calcolo_Parallelo/)
  
- › My contacts
  - [paolo.burgio@unimore.it](mailto:paolo.burgio@unimore.it)
  - <http://hipert.mat.unimore.it/people/paolob/>
  
- › Useful links
  - <http://www.google.com>
  - <http://www.openmp.org>
  - <https://gcc.gnu.org/>
  
- › A "small blog"
  - <http://www.google.com>